

Real-Data Machine Learning for Telemedicine Triage and Outbreak Surveillance

Vaibhav Kadam · Tanvi Gandhi · Omkar Shinolikar · Bylance Technologies

Nexathon II · Paper Presentation · AIKTC · 9 October 2026

Public real-world data only · no synthetic records · released for exact reproduction



IEEE 6-page paper, 13 references



Seed-42 byte-identical reruns



Fail-closed data pipeline

Two decisions, asymmetric errors

1 Patient-level triage

Which callers can be managed remotely, and which need urgent escalation? Errors are asymmetric: a missed high-acuity patient is a safety event, while an unnecessary referral costs access and trust.

Constraint: 15.2% of encounters are high-acuity; uncertain cases must be deferred, not guessed.

2 Population-level outbreak signal

Is this district-week part of a developing wave? In routine public surveillance there is no outbreak label to train on — only a fixed weekly review budget, so every alert displaces another.

Constraint: a top-5% alert budget, and evaluation that must stay descriptive.



Both layers run on ordinary public data — no proprietary EHR, no private sensor feed, no synthetic records.

Three real public datasets — nothing synthetic

13,595

ED analysis encounters

CDC NHAMCS 2019 public microdata (19,481 raw, protocol exclusions applied) [1]

126,712

district-weeks of SARI signals

covid19india public archive, 840 districts, 2020–2023 [2]

69

months of pharmacy sales

Kaggle monthly series, CC BY-NC 4.0 — supplementary forecast task [3]



Missing raw file ⇒ hard FileNotFoundError



Seed 42 — byte-identical reruns




0 synthetic rows anywhere in the pipeline

Every dataset is public and licensed; URLs, retrieval dates and licences are cited in the paper — see Data Availability.

Two modules, one reproducible pipeline

Curago data-science layer: patient-level triage + population-level surveillance



Data: CDC NHAMCS 2019 (public domain) | covid19india archive (static) | Kaggle pharma sales (CC BY-NC, suppl.)

1 Module 1 — patient triage

Calibrated XGBoost (isotonic) risk score → cross-fitted tuned threshold → abstention for uncertain cases → PAC split-conformal recall control at $\alpha = 0.15$.


2 Module 2 — district outbreak signal

IsolationForest anomaly score per district-week → fixed 5% weekly alert budget → top-scoring district-weeks surfaced for review, then described retrospectively.



Both modules recompute from raw CSVs; every table and figure in the paper is a generated artifact — never hand-edited.

Module 1 — discrimination and operating point



Cross-validated ROC, seed 42 — regenerated from `ds/paper_figures.py`.

0.742

primary model CV AUC

vs 0.719 logistic; 95% CI [0.728, 0.751]

17.2%

under-triage, cross-fitted

down from 85.9% at default $t = 0.50$; FPR 0.545

0.829

recall at the chosen point


precision 0.214; thresholds 0.09–0.10 per outer fold

0.113


OOF Brier, isotonic

beats 0.129 for a constant — calibration holds

Split-conformal recall control — marginal vs PAC



Recall distribution over 200 random splits, both rules, $\alpha = 0.15$.



Threshold-coverage trade-off, split means $\pm$ SD.




Marginal rule: mean recall 0.884 (SD 0.033), yet 12.5% of splits fall below the 0.85 target — average-case only.



PAC rule ($\delta = 10\%$): mean recall 0.900, below-target splits cut to 3.5% at precision 0.191.

Module 2 — outbreak surveillance, described not claimed

Outbreak methods vs epidemic windows (descriptive)



Weekly IsolationForest flags across 126,712 district-weeks — regenerated from `ds/paper_figures.py`.

7.0×

Delta-wave enrichment

flag rate inside the wave vs background

2 weeks

median lead over the wave

at the fixed 5% alert budget

25%

of Delta weeks flagged

vs 3.6% of weeks outside it



Descriptive only: no labelled outbreak ground truth exists in routine district data [2]. Enrichment is retrospective description at a fixed budget — not a validated detector.

What we do not claim

- 1 **Guarantees rest on exchangeability** — The marginal rule is average-case (12.5% of splits below target); PAC bounds that share at $\delta = 10\%$, but neither survives distribution shift.
 - 2 **Transfer is unproven** — Trained on US NHAMCS emergency-department data; rural Indian telemedicine casemix and workflow differ. External validation is future work, not a result.
 - 3 **Calibrated, not oracular** — Isotonic Brier 0.113 beats the constant 0.129 — but probabilities are not validated absolute risks; they drive thresholds, not quotations.
 - 4 **Outbreak detection is descriptive** — IsolationForest enrichment is retrospective description at a 5% budget; naive baselines and NEWS2 stay competitive (cutoff 1: precision 0.168, recall 0.669).
-  All ten mandated disclosures are machine-checked in paper/audit.py and pass — full text in Section V of the paper.

Reproducible by construction



Fail-closed data pipeline

A missing or stale raw file raises `FileNotFoundError`. No silent imputation, no synthetic backfill, no cached numbers on disk.



Seed 42, byte-identical

Repeated runs produce identical CSVs and PNGs; published metrics are read back from generated artifacts at build time.



One JSON behind every number

LaTeX, Markdown, DOCX and this deck all resolve metrics from `paper_numbers.json` — zero hand-typed figures.



Automated audits on every build

Numbers, citations, originality, honesty disclosures and IEEE format are machine-checked — all pass.



github.com/vaibhav7087/curago_paper_version · MIT licence · data licences, URLs and retrieval dates in Data Availability

Takeaways

1

Ordinary public data, tuned without test-fold peeking, beats naive baselines on both tasks — AUC 0.742 for triage (cross-fitted under-triage 17.2%) and 7.0× outbreak enrichment — with protocol limits stated up front.

2

Guarantees must be priced: the marginal conformal rule leaves 12.5% of splits below target, and the PAC rule cuts that to 3.5% at $\delta = 10\%$ — publish the failures alongside the mean.

3

Reproducibility is a feature, not a footnote: fail-closed inputs, byte-identical seed-42 reruns, and one JSON behind every number in four output formats.

Next: prospective validation on rural-Indian telemedicine casemix, drift monitoring, threshold re-tuning — SHAP, calibration and forecast artifacts ship with the repository.

References

- [1] National Center for Health Statistics, "National hospital ambulatory medical care survey: 2019 emergency department summary," Centers for Disease Control and Prevention (CDC), 2019, data documentation: https://ftp.cdc.gov/pub/Health_Statistics/NCHS/Dataset_Documentation/NHAMCS/doc19-ed-508.pdf. [Online]. Available: <https://www.cdc.gov/nchs/nhamcs/documentation/about-the-data-2019.html>
- [2] covid19india.com Data Operations Group and incovid19 Team, "Covid-19 india national and district-level time-series archive," Static data archive hosted at <https://data.incovid19.org/>, 2023, retrieved 2026-10-07. Retrospective observed series from 2020-04-26 to 2023-08-22 across 840 state-qualified districts.
- [3] M. Zdravković, "Pharma sales data: Point-of-sale transactional dataset resampled to monthly aggregates," Kaggle, 2019, license: Creative Commons Attribution-NonCommercial 4.0 International (CC BY-NC 4.0). [Online]. Available: <https://www.kaggle.com/datasets/milanzdravkovic/pharma-sales-data>
- [4] Royal College of Physicians, "National Early Warning Score (NEWS) 2: Standardising the Assessment of Acute-Illness Severity in the NHS". London, UK: Royal College of Physicians, 2017, report of a working party.
- [5] G. B. Smith, D. R. Prytherch, P. Meredith, P. E. Schmidt, and P. I. Featherstone, "The ability of the National Early Warning Score (NEWS) to discriminate patients at risk of early death in acute medical admissions," **Resuscitation**, vol. 84, no. 4, pp. 465–470, 2013.
- [6] T. Chen and C. Guestrin, "XGBoost: A scalable tree boosting system," in **Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining (KDD '16)**, 2016, pp. 785–794.
- [7] G. Ke, Q. Meng, T. Finley, T. Wang, W. Chen, W. Ma, Q. Ye, and T.-Y. Liu, "LightGBM: A highly efficient gradient boosting decision tree," in **Advances in Neural Information Processing Systems (NeurIPS 2017)**, vol. 30, 2017, pp. 3146–3154.
- [8] F. Pedregosa, G. Varoquaux, A. Gramfort, V. Michel, B. Thirion, O. Grisel, M. Blondel, P. Prettenhofer, R. Weiss, V. Dubourg, J. Vanderplas, A. Passos, D. Cournapeau, M. Brucher, M. Perrot, and É. Duchesnay, "Scikit-learn: Machine learning in Python," **Journal of Machine Learning Research (JMLR)**, vol. 12, pp. 2825–2830, 2011.
- [9] V. Vovk, A. Gammerman, and G. Shafer, **Algorithmic Learning in a Random World**. New York, NY, USA: Springer Science & Business Media, 2005.
- [10] A. N. Angelopoulos and S. Bates, "A gentle introduction to conformal prediction and distribution-free uncertainty quantification," **arXiv preprint arXiv:2107.07511**, 2021.
- [11] S. M. Lundberg and S.-I. Lee, "A unified approach to interpreting model predictions," in **Advances in Neural Information Processing Systems (NeurIPS 2017)**, vol. 30, 2017, pp. 4765–4774.
- [12] J. Clarke, M. U. G. Kraemer, and et al., "OpenDengue: Data platform for global dengue surveillance," **Scientific Data**, vol. 11, no. 1, p. 324, 2024.
- [13] S. Seabold and J. Perktold, "Statsmodels: Econometric and statistical modeling with Python," in **Proceedings of the 9th Python in Science Conference (SciPy 2010)**, 2010, pp. 92–96.